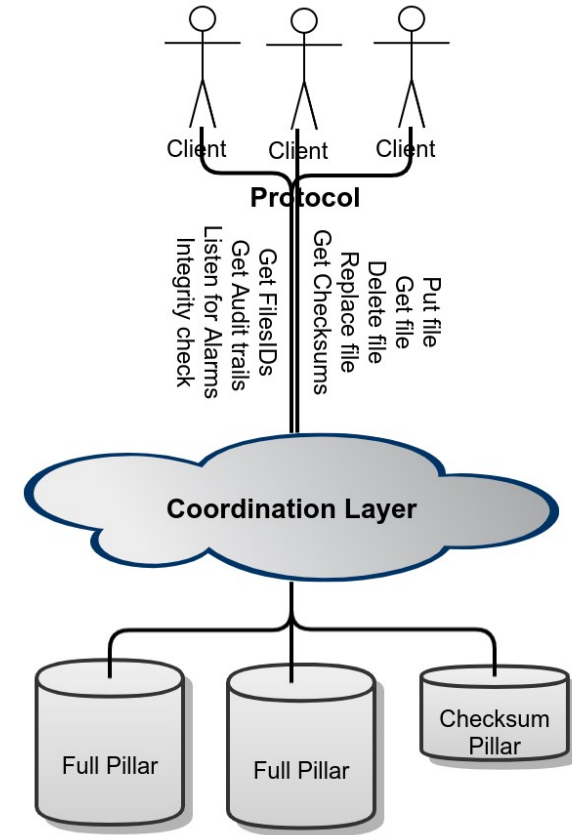


Bitrepository.org

Client training - presentation

What is the bitrepository?

- Distributed storage system
 - Designed for longterm bit-preservation
 - Multiple independent copies
- Uses messagebased asynchronous communication
- Designed to minimize common-mode errors:
 - Implementation (language, OS, etc.)
 - Organizational
 - Storage technology (disk, tapes, optical media)

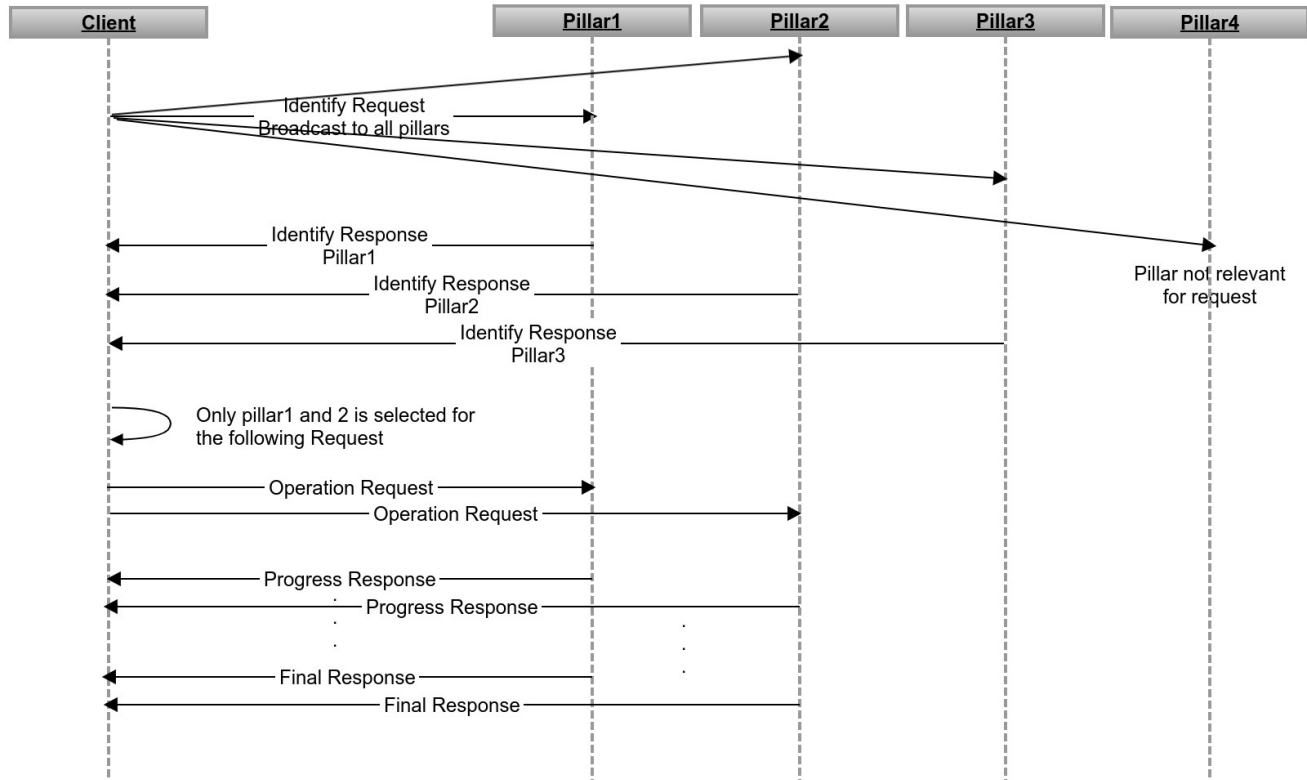


Bitrepository infrastructure

- To function the bitrepository components requires two major bits of infrastructure:
 - Messagebus
 - Transmission of messages between components
 - FileExchange
 - High level abstraction for moving files
 - Initial concept http/https (webdav)
- Enables all other components to be hidden behind firewalls with minimal open ports.

Messageflow

Normal message flow



Conversation

- Operations are handled as a 'conversation'
- A conversation starts with a identification phase
 - The identification phase is for discovery of components that the conversation can continue with
- After a succesful indentification phase, the operation/performing phase can start
- Each conversation have a unique ID

FileExchange

- Common ground to exchange files between components (clients, pillars)
 - Conceptually not limited to a specific protocol
 - Clients responsibility to specify fileexchange url in communication

Security and permission model

- Security is centered around x509 certificates
- RepositorySettings defines the systems full trust.
 - Trusted certificates
 - Permission matrix for each certificate
- Messages are cryptographically signed to:
 - Authenticate a message
 - Authorize an operation

Client library

- One for each protocol primitive
 - PutFileClient, GetChecksumsClient, etc.
- Abstracts most communication complexities away
 - Message signing
 - Messageflow
 - Error senarios
- Interaction with operations are event driven
- A single client can handle multiple concurrent operations

Client library - events

- Good case events

- IDENTIFY_REQUEST_SENT
 - Start of a conversation (identification phase)
- COMPONENT_IDENTIFIED
 - One for each identified component during identification phase
- IDENTIFICATION_COMPLETE
 - Indication that the identification phase is complete
- REQUEST_SENT
 - Start of the operation/performing phase
- PROGRESS
 - A component have delivered a progress response
- COMPONENT_COMPLETE
 - A component have successfully completed its operation
- COMPLETE
 - The entire operation completed successfully

- Bad case events

- COMPONENT_FAILED
 - A component failed its operation
- FAILED
 - The entire operation completed in a failed state
- IDENTIFY_TIMEOUT
 - The identification phase timed out (not all expected contributors responded in time)
- WARNING
 - Something is not quite ok, could for instance be a invalid message.

Client library – factories

- The client library have helper classes to help:
 - Load settings
 - Create MessageBus connection
 - Create the various clients
- Settings are loaded in a static way
 - Provides easy access to settings (the upside)
 - Limits a single JVM to interact with a single bitrepository (the caveat)
- MessageBus connection is worry free:
 - Setup is handled by client factory
 - Closing can be done by MessageBusManager, but should only be done when no more clients are needed in the JVM.

Client library – result paging

- Operations such as GetFileIDs, GetChecksums and GetAuditTrails employs paging to handle long lists of data.
 - Ensures that components scales to millions of files without having to handle all of it at once.
 - One request may therefore only deliver a partial result.
- Paging is done using timestamps
 - GetAuditTrails also by sequence number
- Timestamp paging is inclusive
 - I.e. the last result of the previous operation should be found in the start of the current.
- The component with the smallest pagesize (maxResults) wins
 - i.e. if a pillar is capable of delivering 10.000 results but a client only asks for 1000 pr request, the maximum number of results will be 1000.

Client training - code

- The code for the training can be found at
 - <https://github.com/bitrepository/client-training>

Exercises

- 1) Get client-training project up and running
- 2) Implementation of GetFileIDs
- 3) Implementation of GetFile
- 4) Implementation of PutFile

Exercises

- The exercises will be based on the client-training repository
- The master branch contains runnable code, but only with partial implementation (Operations Get, Put and GetFileIDs are missing, only GetChecksums exists)
- The branch 'solution-proposal' have working implementation of all the implemented operations
- To make the code run, configuration files are needed (RepositorySettings.xml, ReferenceSettings.xml, client-certificate.pem), they should be obtained from a working BitRepository

Exercise 1

- Clone the repository:
<https://github.com/bitrepository/client-training>
 - The project contains boiler plate code for a commandline program, helper classes and a simple implementation for the GetChecksums functionality
- Build the code (master branch)
- Configure the commandline client (product of the build)
- Run a get checksums request on an existing file

Exercise 2

- Implement the GetFileID
 - Start implementation in GetFileIDsAction class
- The implementation should:
 - Request FileIDs from just a single pillar
 - Implement paging of results. Use a page size of 100 results / operation
 - Print list of FileIDs to stdout

Exercise 3

- Implement the GetFile
 - Start implementation in GetAction class
- The implementation should:
 - Request a file from the fastest pillar
 - Use a unique location on the FileExchange (to avoid collisions)
 - Clean up on the FileExchange after the file have been downloaded
 - Not allow to overwrite an existing file on the filesystem

Exercise 4

- Implement the PutFile
 - Start implementation in PutAction class
- The implementation should:
 - Use a unique location on the FileExchange
 - Not upload the file to the FileExchange before identification have completed
 - Cleanup on the FileExchange after the operation